

Healthcare Insurance Analysis

```
In [162... # Let's import the necessary library.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
%matplotlib inline
```

```
In [163... # Let's remove the unnecessary warnings.
import warnings
warnings.filterwarnings("ignore")
```

```
In [164... # Now importing the dataset for the further operation.
cust_details = pd.read_csv("Hospitalisation details.csv")
medical_details = pd.read_csv("Medical Examinations.csv")
cust_name = pd.read_excel("Names.xlsx")
```

```
In [165... cust_details.head()
```

```
Out[165]:
```

| | Customer ID | year | month | date | children | charges | Hospital tier | City tier | State ID |
|---|-------------|------|-------|------|----------|----------|---------------|-----------|----------|
| 0 | Id1 | 1968 | Oct | 12 | 0 | 63770.43 | tier - 1 | tier - 3 | R1013 |
| 1 | Id2 | 1977 | Jun | 8 | 0 | 62592.87 | tier - 2 | tier - 3 | R1013 |
| 2 | Id3 | 1970 | ? | 11 | 3 | 60021.40 | tier - 1 | tier - 1 | R1012 |
| 3 | Id4 | 1991 | Jun | 6 | 1 | 58571.07 | tier - 1 | tier - 3 | R1024 |
| 4 | Id5 | 1989 | Jun | 19 | 0 | 55135.40 | tier - 1 | tier - 2 | R1012 |

```
In [166... cust_details.shape
```

```
Out[166]: (2335, 9)
```

```
In [167... medical_details.head()
```

```
Out[167]:
```

| | Customer ID | BMI | HBA1C | Heart Issues | Any Transplants | Cancer history | NumberOfMajorSurgeries | smoker |
|---|-------------|--------|-------|--------------|-----------------|----------------|------------------------|--------|
| 0 | Id1 | 47.410 | 7.47 | No | No | No | No major surgery | yes |
| 1 | Id2 | 30.360 | 5.77 | No | No | No | No major surgery | yes |
| 2 | Id3 | 34.485 | 11.87 | yes | No | No | 2 | yes |
| 3 | Id4 | 38.095 | 6.05 | No | No | No | No major surgery | yes |
| 4 | Id5 | 35.530 | 5.45 | No | No | No | No major surgery | yes |

```
In [168... medical_details.shape
```

```
Out[168]: (2335, 8)
```

```
In [169... cust_name.head()
```

Out[169]:

| | Customer ID | name |
|---|-------------|-----------------------|
| 0 | Id1 | Hawks, Ms. Kelly |
| 1 | Id2 | Lehner, Mr. Matthew D |
| 2 | Id3 | Lu, Mr. Phil |
| 3 | Id4 | Osborne, Ms. Kelsey |
| 4 | Id5 | Kadala, Ms. Kristyn |

In [170... cust_name.shape

Out[170]: (2335, 2)

Project Task: Week 1

1. Collate the files so that all the information is in one place

In [171...

```
# Now combining the data so that all information could be examine in once go through  
cust_df1 = pd.merge(cust_name, cust_details, on = "Customer ID")  
cust_df1.head()
```

Out[171]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | State ID |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|----------|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | Oct | 12 | 0 | 63770.43 | tier - 1 | tier - 3 | R1013 |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | Jun | 8 | 0 | 62592.87 | tier - 2 | tier - 3 | R1013 |
| 2 | Id3 | Lu, Mr. Phil | 1970 | ? | 11 | 3 | 60021.40 | tier - 1 | tier - 1 | R1012 |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | Jun | 6 | 1 | 58571.07 | tier - 1 | tier - 3 | R1024 |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | Jun | 19 | 0 | 55135.40 | tier - 1 | tier - 2 | R1012 |

In [172...

```
# Now Lets combine the last data set and Complete the all information.  
final_df = pd.merge(cust_df1, medical_details, on = "Customer ID")  
final_df.head()
```

Out[172]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | State ID | BMI | HB |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|----------|--------|----|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | Oct | 12 | 0 | 63770.43 | tier - 1 | tier - 3 | R1013 | 47.410 | |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | Jun | 8 | 0 | 62592.87 | tier - 2 | tier - 3 | R1013 | 30.360 | |
| 2 | Id3 | Lu, Mr. Phil | 1970 | ? | 11 | 3 | 60021.40 | tier - 1 | tier - 1 | R1012 | 34.485 | 1 |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | Jun | 6 | 1 | 58571.07 | tier - 1 | tier - 3 | R1024 | 38.095 | |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | Jun | 19 | 0 | 55135.40 | tier - 1 | tier - 2 | R1012 | 35.530 | |

In [173...]

final_df.shape

Out[173]:

(2335, 17)

2. Check for missing values in the dataset

In [174...]

final_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2335 entries, 0 to 2334
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                          2335 non-null   object
1   name                                 2335 non-null   object
2   year                                 2335 non-null   object
3   month                               2335 non-null   object
4   date                                2335 non-null   int64
5   children                            2335 non-null   int64
6   charges                             2335 non-null   float64
7   Hospital tier                        2335 non-null   object
8   City tier                            2335 non-null   object
9   State ID                            2335 non-null   object
10  BMI                                  2335 non-null   float64
11  HBA1C                               2335 non-null   float64
12  Heart Issues                         2335 non-null   object
13  Any Transplants                     2335 non-null   object
14  Cancer history                      2335 non-null   object
15  NumberOfMajorSurgeries              2335 non-null   object
16  smoker                              2335 non-null   object
dtypes: float64(3), int64(2), object(12)
memory usage: 328.4+ KB
```

In [175...]

final_df.dtypes.value_counts()

Out[175]:

```
object      12
float64      3
int64        2
dtype: int64
```

```
In [176... # Lets check the missing values in the data set.  
final_df.isnull().sum()
```

```
Out[176]: Customer ID          0  
name          0  
year          0  
month         0  
date          0  
children      0  
charges       0  
Hospital tier  0  
City tier      0  
State ID      0  
BMI           0  
HBA1C         0  
Heart Issues  0  
Any Transplants 0  
Cancer history 0  
NumberOfMajorSurgeries 0  
smoker        0  
dtype: int64
```

There is no missing value in the dataset it is clear from the above code, But there is some unusual value that we have to deal.

3. Find the percentage of rows that have trivial value (for example, ?), and delete such rows if they do not contain significant information

```
In [177... trivial_value = final_df[final_df.eq("?").any(1)]  
trivial_value
```

Out[177]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | State ID | BMI |
|------|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|----------|--------|
| 2 | Id3 | Lu, Mr. Phil | 1970 | ? | 11 | 3 | 60021.40 | tier - 1 | tier - 1 | R1012 | 34.485 |
| 169 | Id170 | Torphy, Mr. Bobby | 2000 | Sep | 5 | 1 | 37165.16 | tier - 1 | tier - 3 | ? | 37.620 |
| 559 | Id560 | Pearlman, Mr. Oz | 1994 | Jul | 1 | 3 | 17663.14 | tier - 1 | tier - 3 | R1013 | 23.980 |
| 634 | Id635 | Bruns, Mr. Zachary T | 2004 | Jul | 17 | 0 | 15518.18 | tier - 2 | tier - 3 | R1015 | 25.175 |
| 1285 | Id1286 | Ainsley, Ms. Katie M. | ? | Dec | 12 | 1 | 8547.69 | tier - 2 | tier - 1 | R1013 | 29.370 |
| 1288 | Id1289 | Levine, Ms. Annie J. | ? | Jul | 24 | 0 | 8534.67 | tier - 2 | tier - 3 | R1024 | 24.320 |
| 1792 | Id1793 | Capriolo, Mr. Michael | 1995 | Dec | 1 | 3 | 4827.90 | tier - 1 | tier - 2 | ? | 18.905 |
| 2317 | Id2318 | Gagnon, Ms. Candice M | 1996 | ? | 18 | 0 | 770.38 | tier - 3 | ? | R1012 | 18.820 |
| 2321 | Id2322 | Street, Ms. Holly | 2002 | ? | 19 | 0 | 750.00 | tier - 3 | tier - 1 | R1012 | 21.380 |
| 2323 | Id2324 | Duffy, Ms. Meghan K | 1999 | Dec | 26 | 0 | 700.00 | ? | tier - 3 | R1013 | 22.240 |

In [178... trivial_value.shape

Out[178]: (10, 17)

In [179... *# Percentage of row that have the trivial values*
round(trivial_value.shape[0]/final_df.shape[0]*100, 2)

Out[179]: 0.43

There is total 0.43% of rows contain the trivial values.

In [180... *# Now Lets drop the all row that contain the trivial values in the data set.*
final_df.drop(final_df[final_df.eq("?").any(1)].index, axis=0, inplace=True)

In [181... final_df.shape

Out[181]: (2325, 17)

4. Use the necessary transformation methods to deal with the nominal and ordinal categorical variables in the dataset

```
In [182... # First we will deal with the nominal categorical variable.
```

```
In [183... final_df["Heart Issues"].value_counts()
```

```
Out[183]: No      1405  
yes      920  
Name: Heart Issues, dtype: int64
```

```
In [184... final_df["Any Transplants"].value_counts()
```

```
Out[184]: No      2183  
yes      142  
Name: Any Transplants, dtype: int64
```

```
In [185... final_df["Cancer history"].value_counts()
```

```
Out[185]: No      1934  
Yes      391  
Name: Cancer history, dtype: int64
```

```
In [186... final_df["smoker"].value_counts()
```

```
Out[186]: No      1839  
yes      486  
Name: smoker, dtype: int64
```

```
In [187... # We have some categorical values so first of all we have to transform then by using  
from sklearn.preprocessing import LabelEncoder
```

```
In [188... le = LabelEncoder()
```

```
In [189... final_df["Heart Issues"] = le.fit_transform(final_df["Heart Issues"])  
final_df["Any Transplants"] = le.fit_transform(final_df["Any Transplants"])  
final_df["Cancer history"] = le.fit_transform(final_df["Cancer history"])  
final_df["smoker"] = le.fit_transform(final_df["smoker"])
```

```
In [190... final_df["Heart Issues"].value_counts()
```

```
Out[190]: 0      1405  
1       920  
Name: Heart Issues, dtype: int64
```

```
In [191... final_df.head()
```

Out[191]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | State ID | BMI | HB |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|----------|--------|----|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | Oct | 12 | 0 | 63770.43 | tier - 1 | tier - 3 | R1013 | 47.410 | |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | Jun | 8 | 0 | 62592.87 | tier - 2 | tier - 3 | R1013 | 30.360 | |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | Jun | 6 | 1 | 58571.07 | tier - 1 | tier - 3 | R1024 | 38.095 | |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | Jun | 19 | 0 | 55135.40 | tier - 1 | tier - 2 | R1012 | 35.530 | |
| 5 | Id6 | Baker, Mr. Russell B. | 1962 | Aug | 4 | 0 | 52590.83 | tier - 1 | tier - 3 | R1011 | 32.800 | |

In [192...]

```
# Now we will deal with the ordinal categorical variable.
```

In [193...]

```
def clean_ordinal_variable(val):
    return int(val.replace("tier", "").replace(" ", "").replace("-", ""))
```

In [194...]

```
final_df["Hospital tier"] = final_df["Hospital tier"].map(clean_ordinal_variable)
final_df["City tier"] = final_df["City tier"].map(clean_ordinal_variable)
```

In [195...]

```
final_df["City tier"].value_counts()
```

Out[195]:

```
2    807
3    789
1    729
Name: City tier, dtype: int64
```

In [196...]

```
final_df.head()
```

Out[196]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | State ID | BMI | HB |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|----------|--------|----|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | Oct | 12 | 0 | 63770.43 | 1 | 3 | R1013 | 47.410 | |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | Jun | 8 | 0 | 62592.87 | 2 | 3 | R1013 | 30.360 | |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | Jun | 6 | 1 | 58571.07 | 1 | 3 | R1024 | 38.095 | |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | Jun | 19 | 0 | 55135.40 | 1 | 2 | R1012 | 35.530 | |
| 5 | Id6 | Baker, Mr. Russell B. | 1962 | Aug | 4 | 0 | 52590.83 | 1 | 3 | R1011 | 32.800 | |

5. The dataset has State ID, which has around 16 states. All states are not represented in equal proportions in the data. Creating dummy variables for all regions may also result in too many insignificant predictors. Nevertheless, only R1011, R1012, and R1013 are worth investigating further. Create a suitable strategy to create dummy variables with these restraints.

In [197...]

```
final_df["State ID"].value_counts()
```

Out[197]:

```
R1013    609
R1011    574
R1012    572
R1024    159
R1026     84
R1021     70
R1016     64
R1025     40
R1023     38
R1017     36
R1019     26
R1022     14
R1014     13
R1015     11
R1018      9
R1020      6
Name: State ID, dtype: int64
```

It is clear from the above code some of the state is worth investigator like R1013, R1012, R1011 and R1024.

In [198...]

```
Dummies = pd.get_dummies(final_df["State ID"], prefix= "State_ID")
```

In [199...]

```
Dummies
```


Out[199]:

| | State_ID_R1011 | State_ID_R1012 | State_ID_R1013 | State_ID_R1014 | State_ID_R1015 | State_ID_F |
|------|----------------|----------------|----------------|----------------|----------------|------------|
| 0 | 0 | 0 | 1 | 0 | 0 | |
| 1 | 0 | 0 | 1 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | |
| 4 | 0 | 1 | 0 | 0 | 0 | |
| 5 | 1 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... |
| 2330 | 0 | 0 | 1 | 0 | 0 | |
| 2331 | 0 | 0 | 1 | 0 | 0 | |
| 2332 | 0 | 0 | 1 | 0 | 0 | |
| 2333 | 0 | 0 | 1 | 0 | 0 | |
| 2334 | 0 | 0 | 1 | 0 | 0 | |

2325 rows × 16 columns

In [200... *# Lets take only those state id which play significant role in the data set.*

```
Dummy = Dummies[['State_ID_R1011', 'State_ID_R1012', 'State_ID_R1013']]
Dummy
```

Out[200]:

| | State_ID_R1011 | State_ID_R1012 | State_ID_R1013 |
|------|----------------|----------------|----------------|
| 0 | 0 | 0 | 1 |
| 1 | 0 | 0 | 1 |
| 3 | 0 | 0 | 0 |
| 4 | 0 | 1 | 0 |
| 5 | 1 | 0 | 0 |
| ... | ... | ... | ... |
| 2330 | 0 | 0 | 1 |
| 2331 | 0 | 0 | 1 |
| 2332 | 0 | 0 | 1 |
| 2333 | 0 | 0 | 1 |
| 2334 | 0 | 0 | 1 |

2325 rows × 3 columns

In [201... `final_df = pd.concat([final_df, Dummy], axis=1)`

In [202... `final_df.drop(['State ID'], inplace=True, axis=1)`

In [203... `final_df.head()`

Out[203]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | BMI | HBA1C | H Is |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|--------|-------|------|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | Oct | 12 | 0 | 63770.43 | 1 | 3 | 47.410 | 7.47 | |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | Jun | 8 | 0 | 62592.87 | 2 | 3 | 30.360 | 5.77 | |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | Jun | 6 | 1 | 58571.07 | 1 | 3 | 38.095 | 6.05 | |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | Jun | 19 | 0 | 55135.40 | 1 | 2 | 35.530 | 5.45 | |
| 5 | Id6 | Baker, Mr. Russell B. | 1962 | Aug | 4 | 0 | 52590.83 | 1 | 3 | 32.800 | 6.59 | |

6. The variable NumberOfMajorSurgeries also appears to have string values. Apply a suitable method to clean up this variable.

```
In [204... final_df['NumberOfMajorSurgeries'].value_counts()
```

```
Out[204]: No major surgery    1070
1                      961
2                      272
3                       22
Name: NumberOfMajorSurgeries, dtype: int64
```

The NumberOfMajorSurgeries variable contain string value no major Surgery that mean simpli is 0 surgery so we will replace this value into int value equal to zero.

```
In [205... final_df['NumberOfMajorSurgeries'] = final_df['NumberOfMajorSurgeries'].replace('No
```

```
In [206... final_df['NumberOfMajorSurgeries'] = final_df["NumberOfMajorSurgeries"].astype(int
```

7. Age appears to be a significant factor in this analysis. Calculate the patients' ages based on their dates of birth.

```
In [207... final_df["year"] = pd.to_datetime(final_df["year"], format='%Y').dt.year
final_df["year"]
```

```
Out[207]: 0      1968
          1      1977
          3      1991
          4      1989
          5      1962
          ...
          2330    1998
          2331    1992
          2332    1993
          2333    1992
          2334    1992
Name: year, Length: 2325, dtype: int64
```

```
In [208... final_df["month"] = pd.to_datetime(final_df["month"], format='%b').dt.month
final_df["month"]
```

```
Out[208]: 0      10
          1       6
          3       6
          4       6
          5       8
          ..
          2330     7
          2331     9
          2332     6
          2333    11
          2334     7
Name: month, Length: 2325, dtype: int64
```

```
In [209... final_df['DateInt'] = final_df["year"].astype(str) + final_df["month"].astype(str)
```

```
In [210... final_df['DOB'] = pd.to_datetime(final_df.DateInt, format = "%Y%m%d")
```

```
In [211... final_df.drop(["DateInt"], inplace = True, axis=1)
```

```
In [212... final_df.head()
```

Out[212]:

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | BMI | HBA1C | H Is |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|--------|-------|------|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | 10 | 12 | 0 | 63770.43 | 1 | 3 | 47.410 | 7.47 | |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | 6 | 8 | 0 | 62592.87 | 2 | 3 | 30.360 | 5.77 | |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | 6 | 6 | 1 | 58571.07 | 1 | 3 | 38.095 | 6.05 | |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | 6 | 19 | 0 | 55135.40 | 1 | 2 | 35.530 | 5.45 | |
| 5 | Id6 | Baker, Mr. Russell B. | 1962 | 8 | 4 | 0 | 52590.83 | 1 | 3 | 32.800 | 6.59 | |

```
In [213... import datetime as dt
current_date = dt.datetime.now()
```

```
In [214... final_df['age'] = (((current_date - final_df.DOB).dt.days)/365).astype(int)
```

```
In [215... final_df.head()
```

```
Out[215]:
```

| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | BMI | ... | Heart Issues |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|--------|-----|--------------|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | 10 | 12 | 0 | 63770.43 | 1 | 3 | 47.410 | ... | 0 |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | 6 | 8 | 0 | 62592.87 | 2 | 3 | 30.360 | ... | 0 |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | 6 | 6 | 1 | 58571.07 | 1 | 3 | 38.095 | ... | 0 |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | 6 | 19 | 0 | 55135.40 | 1 | 2 | 35.530 | ... | 0 |
| 5 | Id6 | Baker, Mr. Russell B. | 1962 | 8 | 4 | 0 | 52590.83 | 1 | 3 | 32.800 | ... | 0 |

5 rows × 21 columns

8. The gender of the patient may be an important factor in determining the cost of hospitalization. The salutations in a beneficiary's name can be used to determine their gender. Make a new field for the beneficiary's gender.

```
In [216... def gender(val):
    if "Ms." in val:
        return 0
    else:
        return 1
```

the salutation (Ms.) denote the female and (Mr.) denote the male.

The gender will play the important role to predict the hospitalization cost so for model building we directly denote the gender by int.

Male = 1 & Female = 0

```
In [217... final_df["gender"] = final_df["name"].map(gender)
```

```
In [218... final_df.head()
```

Out[218]:


| | Customer ID | name | year | month | date | children | charges | Hospital tier | City tier | BMI | ... | Transpl |
|---|-------------|-----------------------|------|-------|------|----------|----------|---------------|-----------|--------|-----|---------|
| 0 | Id1 | Hawks, Ms. Kelly | 1968 | 10 | 12 | 0 | 63770.43 | 1 | 3 | 47.410 | ... | |
| 1 | Id2 | Lehner, Mr. Matthew D | 1977 | 6 | 8 | 0 | 62592.87 | 2 | 3 | 30.360 | ... | |
| 3 | Id4 | Osborne, Ms. Kelsey | 1991 | 6 | 6 | 1 | 58571.07 | 1 | 3 | 38.095 | ... | |
| 4 | Id5 | Kadala, Ms. Kristyn | 1989 | 6 | 19 | 0 | 55135.40 | 1 | 2 | 35.530 | ... | |
| 5 | Id6 | Baker, Mr. Russell B. | 1962 | 8 | 4 | 0 | 52590.83 | 1 | 3 | 32.800 | ... | |

5 rows × 22 columns

9. You should also visualize the distribution of costs using a histogram, box and whisker plot, and swarm plot.


In [219...]

```
# Lets make the histogram for the cost distribution.
plt.figure(figsize=(12,5))
sns.histplot(final_df['charges'])
plt.show()
```




In [220...]

```
# Now visualize the cost distribution of the hospitals by box or whisker plot.
plt.boxplot(final_df['charges'])
plt.show()
```



In [221...


```
# Now visualize the cost distribution of the hospitals by swarm plot.
plt.figure(figsize=(12,5))
sns.swarmplot(x='year', y='charges', hue="gender", data=final_df)
plt.xticks(rotation=90)
plt.show()
```



10. State how the distribution is different across gender and tiers of hospitals

In [222...

```
sns.countplot(data = final_df, x='Hospital tier', hue= 'gender')
plt.show()
```



From the above representation it is clear that the number of female in the tier 1 and 3 is half of the male just in tier 2 hospital female is little bit more as compare to male.

11. Create a radar chart to showcase the median hospitalization cost for each tier of hospitals

```
In [223...] print("median cost of tier 1 hospitals:", final_df[final_df["Hospital tier"]==1].cl
print("median cost of tier 2 hospitals:", final_df[final_df["Hospital tier"]==2].cl
print("median cost of tier 3 hospitals:", final_df[final_df["Hospital tier"]==3].cl
```

```
median cost of tier 1 hospitals: 32097.434999999998
median cost of tier 2 hospitals: 7168.76
median cost of tier 3 hospitals: 10676.83
```

```
In [224...] df = pd.DataFrame(dict(r=[32097.43, 7168.76, 10676.83],theta=['tier 1 hospital','t
```

```
In [225...] df
```

```
Out[225]:
```

| | r | theta |
|---|----------|-----------------|
| 0 | 32097.43 | tier 1 hospital |
| 1 | 7168.76 | tier 2 hospital |
| 2 | 10676.83 | tier 3 hospital |

```
In [226...] import plotly.express as px
fig = px.line_polar(df, r='r', theta='theta', line_close=True)
fig.update_traces(fill='toself')
fig.show()
```

12. Create a frequency table and a stacked bar chart to visualize the count of people in the different tiers of cities and hospitals

```
In [227... # Frequency table for count of the people according to the tier of city and hospital
final_df["Hospital tier"].value_counts()
```

Out[227]:

| | |
|---|------|
| 2 | 1334 |
| 3 | 691 |
| 1 | 300 |

Name: Hospital tier, dtype: int64

```
In [228... city_freq = final_df["City tier"].value_counts().rename_axis('City&hospital_tier')
```

```
In [229... hospital_freq = final_df["Hospital tier"].value_counts().rename_axis('City&hospital_tier')
```


```
In [230... df = pd.merge(city_freq, hospital_freq, on = 'City&hospital_tier')
```

```
In [231... df
```

Out[231]:

| | City&hospital_tier | city_counts | hospital_counts |
|---|--------------------|-------------|-----------------|
| 0 | 2 | 807 | 1334 |
| 1 | 3 | 789 | 691 |
| 2 | 1 | 729 | 300 |


```
In [232... plt.bar(df["City&hospital_tier"], df["city_counts"], color='r')
plt.bar(df["City&hospital_tier"], df["hospital_counts"], bottom=df["city_counts"],
plt.show()
```



13. Test the following null hypotheses:

- The average hospitalization costs for the three types of hospitals are not significantly different
- The average hospitalization costs for the three types of cities are not significantly different
- The average hospitalization cost for smokers is not significantly different from the average cost for nonsmokers
- Smoking and heart issues are independent

```
In [233... from scipy.stats import ttest_1samp
```

```
In [234... # a. The average hospitalization costs for the three types of hospitals are not sig
print("median cost of tier 1 hospitals:", final_df[final_df["Hospital tier"]==1].cl
print("median cost of tier 2 hospitals:", final_df[final_df["Hospital tier"]==2].cl
print("median cost of tier 3 hospitals:", final_df[final_df["Hospital tier"]==3].cl
```

```
median cost of tier 1 hospitals: 32097.434999999998
median cost of tier 2 hospitals: 7168.76
median cost of tier 3 hospitals: 10676.83
```

Interpretation

H0: the distributions of all samples are equal. || H1: the distributions of one or more samples are not equal

```
In [235... from scipy.stats import friedmanchisquare
data1 = [32097.43]
data2 = [7168.76]
data3 = [10676.83]
stat, p = friedmanchisquare(data1, data2, data3)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=2.000, p=0.368
Probably the same distribution

```
In [236... # b. The average hospitalization costs for the three types of cities are not signi
print("median cost of tier 1 city:", final_df[final_df["City tier"]==1].charges.med
print("median cost of tier 2 city:", final_df[final_df["City tier"]==2].charges.med
print("median cost of tier 3 city:", final_df[final_df["City tier"]==3].charges.med
```

median cost of tier 1 city: 10027.15
median cost of tier 2 city: 8968.33
median cost of tier 3 city: 9880.07

```
In [237... data1 = [10027.15]
data2 = [8968.33]
data3 = [9880.07]
stat, p = friedmanchisquare(data1, data2, data3)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=2.000, p=0.368
Probably the same distribution

```
In [238... # c. The average hospitalization cost for smokers is not significantly different fr
print("median cost of smoker:", final_df[final_df["smoker"]==1].charges.median())
print("median cost of non smoker:", final_df[final_df["smoker"]==0].charges.median
```

median cost of smoker: 34125.475
median cost of non smoker: 7537.16

```
In [239... from scipy.stats import kruskal
data1 = [34125.475]
data2 = [7537.16]
stat, p = kruskal(data1, data2)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=1.000, p=0.317
Probably the same distribution

Interpretation

H0: the two samples are independent. H1: there is a dependency between the samples.

```
In [240... # d. Smoking and heart issues are independent
from scipy.stats import chi2_contingency
table = [[final_df["Heart Issues"].value_counts(), [final_df["smoker"].value_count
```

```
stat, p, dof, expected = chi2_contingency(table)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably independent')
else:
    print('Probably dependent')
```

stat=191.145, p=0.000
Probably dependent

Project Task: Week 2

1. Examine the correlation between predictors to identify highly correlated predictors. Use a heatmap to visualize this.

In [241...

```
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2325 entries, 0 to 2334
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                          2325 non-null   object
1   name                                 2325 non-null   object
2   year                                 2325 non-null   int64
3   month                               2325 non-null   int64
4   date                                2325 non-null   int64
5   children                            2325 non-null   int64
6   charges                             2325 non-null   float64
7   Hospital tier                        2325 non-null   int64
8   City tier                            2325 non-null   int64
9   BMI                                 2325 non-null   float64
10  HBA1C                               2325 non-null   float64
11  Heart Issues                         2325 non-null   int32
12  Any Transplants                     2325 non-null   int32
13  Cancer history                      2325 non-null   int32
14  NumberOfMajorSurgeries              2325 non-null   int32
15  smoker                              2325 non-null   int32
16  State_ID_R1011                      2325 non-null   uint8
17  State_ID_R1012                      2325 non-null   uint8
18  State_ID_R1013                      2325 non-null   uint8
19  DOB                                 2325 non-null   datetime64[ns]
20  age                                 2325 non-null   int32
21  gender                              2325 non-null   int64
dtypes: datetime64[ns](1), float64(3), int32(6), int64(7), object(2), uint8(3)
memory usage: 315.6+ KB
```

In [242...

```
# In the data frame some of the column are not usable to model building so Lets find
#then indentify the highly corelated predictor.
final_df.drop(["Customer ID", 'name', 'year', 'month', 'date', 'DOB'], inplace=True)
```

In [243...

```
final_df.shape
```

Out[243]:

```
(2325, 16)
```

In [244...

```
final_df.head()
```

Out[244]:

| | children | charges | Hospital tier | City tier | BMI | HBA1C | Heart Issues | Any Transplants | Cancer history | NumberOfMajorSurgeries |
|---|----------|----------|---------------|-----------|--------|-------|--------------|-----------------|----------------|------------------------|
| 0 | 0 | 63770.43 | 1 | 3 | 47.410 | 7.47 | 0 | 0 | 0 | |
| 1 | 0 | 62592.87 | 2 | 3 | 30.360 | 5.77 | 0 | 0 | 0 | |
| 3 | 1 | 58571.07 | 1 | 3 | 38.095 | 6.05 | 0 | 0 | 0 | |
| 4 | 0 | 55135.40 | 1 | 2 | 35.530 | 5.45 | 0 | 0 | 0 | |
| 5 | 0 | 52590.83 | 1 | 3 | 32.800 | 6.59 | 0 | 0 | 0 | |

In [245...


```
corr = final_df.corr()  
corr
```

Out[245]:

| | children | charges | Hospital tier | City tier | BMI | HBA1C | Heart Issues | Any Transplants | Cancer history | NumberOfMajorSurgeries | smoker | State_ID_R1011 | State_ID_R1012 | State_ID_R1013 | age | gender |
|------------------------|-----------|-----------|---------------|-----------|-----------|-----------|--------------|-----------------|----------------|------------------------|-----------|----------------|----------------|----------------|-----------|-----------|
| children | 1.000000 | 0.055901 | -0.052438 | -0.015760 | -0.005339 | -0.101379 | 0.023984 | -0.142040 | -0.027880 | -0.113161 | 0.017713 | 0.011666 | 0.005247 | -0.013834 | -0.005457 | 0.011205 |
| charges | 0.055901 | 1.000000 | -0.446687 | 0.035300 | 0.346730 | 0.139697 | 0.049299 | -0.127028 | -0.022522 | 0.053308 | 0.838462 | 0.286956 | -0.074636 | -0.150634 | 0.304395 | 0.034069 |
| Hospital tier | -0.052438 | -0.446687 | 1.000000 | -0.039755 | -0.104771 | 0.057855 | 0.053376 | 0.011729 | -0.021429 | 0.033230 | -0.474077 | -0.114685 | 0.020272 | 0.002455 | 0.133771 | 0.041261 |
| City tier | -0.015760 | 0.035300 | -0.039755 | 1.000000 | 0.038123 | -0.005404 | 0.023152 | 0.002970 | -0.018639 | 0.027937 | 0.032034 | 0.036049 | -0.018253 | 0.002766 | -0.008070 | 0.054073 |
| BMI | -0.005339 | 0.346730 | -0.104771 | 0.038123 | 1.000000 | -0.006920 | 0.017129 | 0.015893 | -0.020235 | 0.018851 | 0.107126 | 0.115671 | 0.017939 | -0.208744 | 0.049260 | 0.079930 |
| HBA1C | -0.101379 | 0.139697 | 0.057855 | -0.005404 | -0.006920 | 1.000000 | 0.007699 | -0.159855 | -0.170921 | -0.091594 | 0.007257 | 0.015525 | -0.019513 | 0.033453 | 0.460558 | -0.027339 |
| Heart Issues | 0.023984 | 0.049299 | 0.053376 | 0.023152 | 0.017129 | 0.007699 | 1.000000 | -0.140261 | 0.111195 | 0.206141 | -0.007195 | 0.005895 | 0.021771 | -0.027961 | 0.192271 | 0.010271 |
| Any Transplants | -0.142040 | -0.127028 | 0.011729 | 0.002970 | 0.015893 | -0.159855 | -0.140261 | 1.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 |
| Cancer history | -0.027880 | -0.022522 | -0.021429 | -0.018639 | -0.020235 | -0.170921 | 0.111195 | -0.000000 | 1.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 |
| NumberOfMajorSurgeries | -0.113161 | 0.053308 | 0.033230 | 0.027937 | 0.018851 | -0.091594 | 0.206141 | -0.000000 | -0.000000 | 1.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 |
| smoker | 0.017713 | 0.838462 | -0.474077 | 0.032034 | 0.107126 | 0.007257 | -0.007195 | -0.000000 | -0.000000 | -0.000000 | 1.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 |
| State_ID_R1011 | 0.011666 | 0.286956 | -0.114685 | 0.036049 | 0.115671 | 0.015525 | 0.005895 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | 1.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 |
| State_ID_R1012 | 0.005247 | -0.074636 | 0.020272 | -0.018253 | 0.017939 | -0.019513 | 0.021771 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | 1.000000 | -0.000000 | -0.000000 | -0.000000 |
| State_ID_R1013 | -0.013834 | -0.150634 | 0.002455 | 0.002766 | -0.208744 | 0.033453 | -0.027961 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | 1.000000 | -0.000000 | -0.000000 |
| age | -0.005457 | 0.304395 | 0.133771 | -0.008070 | 0.049260 | 0.460558 | 0.192271 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | 1.000000 | -0.000000 |
| gender | 0.011205 | 0.034069 | 0.041261 | 0.054073 | 0.079930 | -0.027339 | 0.010271 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | -0.000000 | 1.000000 |

In [246...

```
plt.figure(figsize=(15,10))  
sns.heatmap(corr, annot=True, linewidth=.5, cmap="crest")  
plt.show()
```



From the above correlation its clear that somker variable is highly corealted to the output variable.

2. Develop and evaluate the final model using regression with a stochastic gradient descent optimizer. Also, ensure that you apply all the following suggestions:

Note:

- Perform the stratified 5-fold cross-validation technique for model building and validation • Use standardization and hyperparameter tuning effectively • Use sklearn-pipelines • Use appropriate regularization techniques to address the bias-variance trade-off

a. Create five folds in the data, and introduce a variable to identify the folds

b. For each fold, run a for loop and ensure that 80 percent of the data is used to train the model and the remaining 20 percent is used to validate it in each iteration

c. Develop five distinct models and five distinct validation scores (root mean squared error values)

d. Determine the variable importance scores, and identify the redundant variables

```
In [247... # Lets first separate the input and output data.
x = final_df.drop(["charges"], axis=1)
y = final_df[['charges']]
```

```
In [248... # Lets split the data set into the training and testing data.
from sklearn.model_selection import train_test_split
```

```
In [249... x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=.20, random_st
```

```
In [250... # Now standardize the data.
from sklearn.preprocessing import StandardScaler
```

```
In [251... sc = StandardScaler()
```

```
In [252... x_train = sc.fit_transform(x_train)
x_test = sc.fit_transform(x_test)
```

```
In [253... from sklearn.linear_model import SGDRegressor
```

```
In [254... from sklearn.model_selection import GridSearchCV

params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2,0.3,0.4,0.5,
                    0.6,0.7,0.8,0.9,1.0,2.0,3.0,4.0,5.0,6.0,7.0,8.0,
                    9.0,10.0,20,50,100,500,1000],
          'penalty': ['l2', 'l1', 'elasticnet']}

sgd = SGDRegressor()

# Cross Validation
folds = 5
model_cv = GridSearchCV(estimator = sgd,
                        param_grid = params,
                        scoring = 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score = True,
                        verbose = 1)
model_cv.fit(x_train,y_train)
```

```
Out[254]: Fitting 5 folds for each of 84 candidates, totalling 420 fits
GridSearchCV(cv=5, estimator=SGDRegressor(),
             param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
                                   0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                                   4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
                                   100, 500, 1000],
                         'penalty': ['l2', 'l1', 'elasticnet']}},
             return_train_score=True, scoring='neg_mean_absolute_error',
             verbose=1)
```

```
In [255... model_cv.best_params_
```

```
Out[255]: {'alpha': 50, 'penalty': 'l1'}
```

```
In [256... sgd = SGDRegressor(alpha= 100, penalty= 'l1')
```

```
In [257... sgd.fit(x_train, y_train)
```

```
Out[257]: SGDRegressor(alpha=100, penalty='l1')
```

```
In [258... sgd.score(x_test, y_test)
```

Out[258]: 0.8574058061920549

```
In [259... y_pred = sgd.predict(x_test)
```

```
In [260... from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [261... sgd_mae = mean_absolute_error(y_test, y_pred)
sgd_mse = mean_squared_error(y_test, y_pred)
sgd_rmse = np.sqrt(sgd_mse)
```

```
In [262... print("MAE:", sgd_mae)
print("MSE:", sgd_mse)
print("RMSE:", sgd_rmse)
```

MAE: 3145.3196499430524
MSE: 23985187.279824603
RMSE: 4897.4674353000655

```
In [263... # d. Determine the variable importance scores, and identify the redundant variables
importance = sgd.coef_
```

```
In [264... pd.DataFrame(importance, index = x.columns, columns=['Feature_imp'])
```

Out[264]:

| | Feature_imp |
|------------------------|--------------|
| children | 365.627972 |
| Hospital tier | -1080.163416 |
| City tier | 0.000000 |
| BMI | 2669.591053 |
| HBA1C | 41.799507 |
| Heart Issues | 0.000000 |
| Any Transplants | 0.000000 |
| Cancer history | 47.258781 |
| NumberOfMajorSurgeries | 0.000000 |
| smoker | 8741.236701 |
| State_ID_R1011 | -150.534783 |
| State_ID_R1012 | 0.000000 |
| State_ID_R1013 | -376.626750 |
| age | 3383.403770 |
| gender | 0.000000 |

| | Feature_imp |
|------------------------|--------------|
| children | 365.627972 |
| Hospital tier | -1080.163416 |
| City tier | 0.000000 |
| BMI | 2669.591053 |
| HBA1C | 41.799507 |
| Heart Issues | 0.000000 |
| Any Transplants | 0.000000 |
| Cancer history | 47.258781 |
| NumberOfMajorSurgeries | 0.000000 |
| smoker | 8741.236701 |
| State_ID_R1011 | -150.534783 |
| State_ID_R1012 | 0.000000 |
| State_ID_R1013 | -376.626750 |
| age | 3383.403770 |
| gender | 0.000000 |

3. Use random forest and extreme gradient boosting for cost prediction, share your crossvalidation results, and calculate the variable importance scores

random forest

```
In [265... from sklearn.ensemble import RandomForestRegressor
```

```
In [266... # Instantiate model with 1000 decision trees
rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)

# Train the model on training data
rf.fit(x_train, y_train)
```

```
Out[266]: RandomForestRegressor(n_estimators=1000, random_state=42)
```

```
In [267... score = rf.score(x_test,y_test)
score
```

```
Out[267]: 0.9222696338245824
```

```
In [268... y_pred = rf.predict(x_test)
```

```
In [269... rf_mae = mean_absolute_error(y_test, y_pred)
```

```
In [270... rf_mae
```

```
Out[270]: 1870.3529629462323
```

extreme gradient boosting

```
In [271... from sklearn.ensemble import GradientBoostingRegressor
```

```
In [272... # Instantiate model with 1000 decision trees
gbr = GradientBoostingRegressor(n_estimators = 1000, random_state = 42)

# Train the model on training data
gbr.fit(x_train, y_train)
```

```
Out[272]: GradientBoostingRegressor(n_estimators=1000, random_state=42)
```

```
In [273... score = gbr.score(x_test,y_test)
score
```

```
Out[273]: 0.9042734212625119
```

```
In [274... y_pred = gbr.predict(x_test)
```

```
In [275... gbr_mae = mean_absolute_error(y_test, y_pred)
gbr_mae
```

```
Out[275]: 2375.8700944163274
```

4. Case scenario:

Estimate the cost of hospitalization for Christopher, Ms. Jayna (her date of birth is 12/28/1988, height is 170 cm, and weight is 85 kgs). She lives in a tier-1 city and her state's State ID is R1011. She lives with her partner and two children. She was found to be nondiabetic (HbA1c = 5.8). She smokes but is otherwise healthy. She has had no transplants or major

surgeries. Her father died of lung cancer. Hospitalization costs will be estimated using tier-1 hospitals.

```
In [276... # First we need to calculate the age of the person.
date = "19881228"
date1 = datetime.strptime(date,"%Y%m%d")
date1
```

```
Out[276]: datetime.datetime(1988, 12, 28, 0, 0)
```

```
In [277... current_date = datetime.now()
current_date
```

```
Out[277]: datetime.datetime(2023, 3, 1, 22, 56, 36, 990464)
```

```
In [278... age = str((current_date - date1)/365)
```

```
In [279... print("Age=",age[:2])
```

```
Age= 34
```

```
In [280... # now with the help of height and weight we will calculate the BMI.
height_m = 170/100
height_sq = height_m*height_m
BMI = 85/height_sq
np.round(BMI,2)
```

```
Out[280]: 29.41
```

```
In [281... # Now Lets gen
list = [[2,1,1,24.41,5.8,0,0,0,0,1,1,0,0,34,0]]
```

```
In [282... df = pd.DataFrame(list, columns = ['children', 'Hospital tier', 'City tier', 'BMI',
                                   'Cancer history', 'NumberOfMajorSurgeries', 'smoker',
                                   'State_ID_R1013', 'age', 'gender'] )
df
```

```
Out[282]:
```

| | children | Hospital tier | City tier | BMI | HBA1C | Heart Issues | Any Transplants | Cancer history | NumberOfMajorSurgeries |
|---|----------|---------------|-----------|-------|-------|--------------|-----------------|----------------|------------------------|
| 0 | 2 | 1 | 1 | 24.41 | 5.8 | 0 | 0 | 0 | 0 |

5. Find the predicted hospitalization cost using all models. The predicted value should be the mean of the five models' predicted values.

```
In [283... Hospital_cost = []
```

```
In [284... # Now Lets predict the hospitalization cost through SGDRegressor
Cost1 = sgd.predict(df)
Hospital_cost.append(Cost1)
```

```
In [285... # Now Lets predict the hospitalization cost through Random Forest
Cost2 = rf.predict(df)
Hospital_cost.append(Cost2)
```

```
In [286... # Now Lets predict the hospitalization cost throug Extreme gradient Booster  
Cost3 = gbr.predict(df)  
Hospital_cost.append(Cost3)
```

```
In [287... avg_cost = np.mean(Hospital_cost)  
avg_cost
```

```
Out[287]: 103919.51224697009
```

So in the new case the avg predicted hospitalization cost is 103919.51

```
In [ ]:
```